

# zillow\_MWS

October 13, 2017

## 1 Zillow prize data analysis report

```
In [1]: from datetime import datetime
        d = datetime.now().date()
        t = datetime.now().strftime('%H:%M:%S')
        print("This report was last updated on", d, "at", t)
```

This report was last updated on 2017-10-12 at 11:52:31

### 1.1 Introduction

The [Zillow Prize](#) is a [Kaggle competition](#) that aims to inspire data scientists around the world to improve the accuracy of the Zillow "Zestimate" statistical and machine learning models.

My goal is to compete for the Zillow prize and write up my results.

### 1.2 Methods

#### 1.2.1 Data

The data were obtained from [Kaggle website](#) and consist of the following files: - properties\_2016.csv.zip - properties\_2017.csv.zip - sample\_submission.csv - train\_2016\_v2.csv.zip - train\_2017.csv.zip - zillow\_data\_dictionary.xlsx The zillow\_data\_dictionary.xlsx is a code book that explains the data. This data will be made available on [figshare](#) to provide an additional source if the [Kaggle site data](#) become unavailable.

#### 1.2.2 Analysis

Data analysis was done in Jupyter Notebook (Pérez and Granger 2007)[?] Integrated Development Environment using the Python language (Pérez, Granger, and Hunter 2011)[?] and a number of software packages:

- NumPy (van der Walt, Colbert, and Varoquaux 2011)[?]
- pandas (McKinney 2010)[?]
- scikit-learn (Pedregosa et al. 2011)[?]

### 1.2.3 Visualization

The following packages were used to visualize the data:

- Matplotlib (Hunter 2007)[?]
- Seaborn (Waskom et al. 2014)[?]
- r-ggplot2
- r-cowplot

The use of R code and packages in a Python environment is possible through the use of the Rpy2 package.

### 1.2.4 Prediction

Machine learning prediction was done using the following packages:

- scikit-learn (Pedregosa et al. 2011)[?]
- xgboost
- r-caret

### 1.2.5 Reproducibility

Reproducibility is extremely important in scientific research yet many examples of problematic studies exist in the literature (Couzin-Frankel 2010)[?].

The names and versions of each package used herein are listed in the accompanying `env.yml` file in the `config` folder. The computational environment used to analyze the data can be recreated using this `env.yml` file and the [conda package and environment manager](#) available as part of the [Anaconda distribution of Python](#).

Additionally, details on how to setup a Docker image capable of running the analysis is included in the `README.md` file in the `config` folder.

The code in the form of a jupyter notebook (`01_zillow_MWS.ipynb`) or Python script (`01_zillow_MWS.py`), can also be run on the Kaggle website (this requires logging in with a username and password).

More information on the details of how this project was created and the computational environment was configured can be found in the accompanying `README.md` file.

This Python 3 environment comes with many helpful analytics libraries installed It is defined by the kaggle/python docker image: <https://github.com/kaggle/docker-python> (a modified version of this docker image will be made available as part of my project to ensure reproducibility). For example, here's several helpful packages to load in

## 1.3 Results

### 1.3.1 Import Libraries and Data

Input data files are available in the `"../input/"` directory.

Any results I write to the current directory are saved as output.

```
In [2]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # data visualization
import datetime as dt
import seaborn as sns
import xgboost as xgb
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestRegressor
from sklearn import ensemble
%matplotlib inline
### Seaborn style
sns.set_style("whitegrid")
```

/Users/marskar/anaconda3/lib/python3.6/site-packages/sklearn/cross\_validation.py:41:  
 DeprecationWarning: This module was deprecated in version 0.18 in favor of the  
 model\_selection module into which all the refactored classes and functions are moved.  
 Also note that the interface of the new CV iterators are different from that of this  
 module. This module will be removed in 0.20.  
 "This module will be removed in 0.20.", DeprecationWarning)


```
In [3]: prop = pd.read_csv("../input/properties_2016.csv")
prop.shape
```

/Users/marskar/anaconda3/lib/python3.6/site-  
 packages/IPython/core/interactiveshell.py:2728: DtypeWarning: Columns (22,32,34,49,55)  
 have mixed types. Specify dtype option on import or set low\_memory=False.  
 interactivity=interactivity, compiler=compiler, result=result)

```
In [4]: ### ... check for NaNs
nan = prop.isnull().sum()/len(prop)*100
nan
```

```
In [5]: ### Plotting NaN counts
nan_sorted = nan.sort_values(ascending=False).to_frame().reset_index()
nan_sorted.columns = ['Column', 'Number of NaNs']
```

```
In [6]: fig, ax = plt.subplots(figsize=(12, 25))
sns.barplot(x="Number of NaNs", y="Column", data=nan_sorted, color='Blue', ax=ax)
ax.set(xlabel="Number of NaNs", ylabel="", title="Total Number of NaNs in each column")
plt.show()
```



There are several columns which have a very high proportion of missing values. It may be worth analysing these more closely.

## Feature Importance by Random Forest

```
In [7]: train = pd.read_csv("../input/train_2016_v2.csv", parse_dates=["transactiondate"])
        train.shape
```

```
In [8]: train['transaction_month'] = pd.DatetimeIndex(train['transactiondate']).month
        train.sort_values('transaction_month', axis=0, ascending=True, inplace=True)
```

## Feature Importance

```
In [9]: #fill NaN values with -1 and encode object columns
        for x in prop.columns:
            prop[x] = prop[x].fillna(-1)

        #many more parcelids in properties file, merge with training file
        train = pd.merge(train, prop, on='parcelid', how='left')

In [10]: for c in train[['transactiondate', 'hashottuborspa', 'propertycountylandusecode',
                        'propertyzoningdesc', 'fireplaceflag', 'taxdelinquencyflag']]:
            label = LabelEncoder()
            label.fit(list(train[c].values))
            train[c] = label.transform(list(train[c].values))

            x_train = train.drop(['parcelid', 'logerror', 'transactiondate'], axis=1)
            y_train = train['logerror']

In [11]: rf = RandomForestRegressor(n_estimators=30, max_features=None)
        rf.fit(x_train, y_train)


In [12]: rf_importance = rf.feature_importances_
        rf_importance_df = pd.DataFrame()
        rf_importance_df['features'] = x_train.columns
        rf_importance_df['importance'] = rf_importance
        print(rf_importance_df.head())
```

	features	importance
0	transaction_month	0.041078
1	airconditioningtypeid	0.005809
2	architecturalstyletypeid	0.000262
3	basementsqft	0.000238
4	bathroomcnt	0.007725

```
In [13]: rf_importance_df.sort_values('importance', axis=0, inplace=True, ascending=False)
        print(rf_importance_df.head())
```

	features	importance
50	structuretaxvaluedollarcnt	0.084278
24	latitude	0.078829
54	taxamount	0.076839
25	longitude	0.072749
26	lotsizesquarefeet	0.071840

```
In [14]: fig, ax = plt.subplots(figsize=(6, 12.5))
sns.barplot(x="importance", y="features", data=rf_importance_df, color='Blue', ax=ax)
ax.set(xlabel="Importance", ylabel="Feature Name", title="Feature Importances")
plt.show()
```



```
In [15]: etr = ensemble.ExtraTreesRegressor(n_estimators=25, max_depth=30, max_features=0.3,
      n_jobs=-1, random_state=0)
      etr.fit(x_train, y_train)
```

```
In [16]: etr_importance = etr.feature_importances_
      etr_importance_df = pd.DataFrame()
      etr_importance_df['features'] = x_train.columns
      etr_importance_df['importance'] = etr_importance
      print(etr_importance_df.head())
```

	features	importance
0	transaction_month	0.052949
1	airconditioningtypeid	0.015291
2	architecturalstyletypeid	0.000204
3	basementsqft	0.000212
4	bathroomcnt	0.014486


```
In [17]: etr_importance_df.sort_values('importance', axis=0, inplace=True, ascending=False)

      print(etr_importance_df.head())
```

	features	importance
50	structuretaxvaluedollarcnt	0.060509
54	taxamount	0.059521
53	landtaxvaluedollarcnt	0.056897
51	taxvaluedollarcnt	0.056176
26	lotssizesquarefeet	0.054112

```
In [18]: fig, ax = plt.subplots(figsize=(6, 12.5))
      sns.barplot(x="importance", y="features", data=etr_importance_df, color='Blue', ax=ax)
      ax.set(xlabel="Importance", ylabel="Feature Name", title="Feature Importances")
      plt.show()
```






```

In [19]: xgb_params = {
    'eta': 0.05,
    'max_depth': 8,
    'subsample': 0.7,
    'colsample_bytree': 0.7,
    'objective': 'reg:linear',
    'silent': 1,
    'seed': 0
}
dtrain = xgb.DMatrix(x_train, y_train, feature_names=x_train.columns.values)
model = xgb.train(dict(xgb_params, silent=0), dtrain, num_boost_round=50)

In [20]: # plot the important features #
fig, ax = plt.subplots(figsize=(6,9))
xgb.plot_importance(model, height=0.8, grid = False, color="blue", ax=ax)
ax.xaxis.grid()
plt.show()

```



```

In [ ]: import matplotlib.pyplot as plt
import numpy as np

# Simple data to display in various forms
x = np.linspace(0, 2 * np.pi, 400)
y = np.sin(x ** 2)

plt.close('all')

# Just a figure and one subplot
f, ax = plt.subplots()
ax.plot(x, y)
ax.set_title('Simple plot')

# Two subplots, the axes array is 1-d
f, axarr = plt.subplots(2, sharex=True)
axarr[0].plot(x, y)
axarr[0].set_title('Sharing X axis')
axarr[1].scatter(x, y)

# Two subplots, unpack the axes array immediately
f, (ax1, ax2) = plt.subplots(1, 2, sharey=True)
ax1.plot(x, y)
ax1.set_title('Sharing Y axis')
ax2.scatter(x, y)

# Three subplots sharing both x/y axes
f, (ax1, ax2, ax3) = plt.subplots(3, sharex=True, sharey=True)
ax1.plot(x, y)
ax1.set_title('Sharing both axes')
ax2.scatter(x, y)
ax3.scatter(x, 2 * y ** 2 - 1, color='r')
# Fine-tune figure; make subplots close to each other and hide x ticks for
# all but bottom plot.
f.subplots_adjust(hspace=0)
plt.setp([a.get_xticklabels() for a in f.axes[:-1]], visible=False)

# row and column sharing
f, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, sharex='col', sharey='row')
ax1.plot(x, y)
ax1.set_title('Sharing x per column, y per row')
ax2.scatter(x, y)
ax3.scatter(x, 2 * y ** 2 - 1, color='r')
ax4.plot(x, 2 * y ** 2 - 1, color='r')

# Four axes, returned as a 2-d array
f, axarr = plt.subplots(2, 2)
axarr[0, 0].plot(x, y)
axarr[0, 0].set_title('Axis [0,0]')
axarr[0, 1].scatter(x, y)
axarr[0, 1].set_title('Axis [0,1]')
axarr[1, 0].plot(x, y ** 2)
axarr[1, 0].set_title('Axis [1,0]')
axarr[1, 1].scatter(x, y ** 2)
axarr[1, 1].set_title('Axis [1,1]')
# Fine-tune figure; hide x ticks for top plots and y ticks for right plots
plt.setp([a.get_xticklabels() for a in axarr[0, :]], visible=False)
plt.setp([a.get_yticklabels() for a in axarr[:, 1]], visible=False)

# Four polar axes
f, axarr = plt.subplots(2, 2, subplot_kw=dict(projection='polar'))
axarr[0, 0].plot(x, y)
axarr[0, 0].set_title('Axis [0,0]')
axarr[0, 1].scatter(x, y)
axarr[0, 1].set_title('Axis [0,1]')

```

```

axarr[1, 0].plot(x, y ** 2)
axarr[1, 0].set_title('Axis [1,0]')
axarr[1, 1].scatter(x, y ** 2)
axarr[1, 1].set_title('Axis [1,1]')
# Fine-tune figure; make subplots farther from each other.
f.subplots_adjust(hspace=0.3)

plt.show()

```

## 1.4 Conclusions

In Progress

## 1.5 Bibliography

Couzin-Frankel, J. 2010. “Cancer Research. As Questions Grow, Duke Halts Trials, Launches Investigation.” *Science* 329 (5992): 614–15.

Hunter, J. D. 2007. “Matplotlib: A 2D Graphics Environment.” *Computing In Science & Engineering* 9 (3): 90–95.

McKinney, W. 2010. “Data Structures for Statistical Computing in Python.” In *Proceedings of the 9th Python in Science Conference*, edited by S. J. van der Walt and K. J. Millman. Austin, Texas.

Pedregosa, Fabian, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, et al. 2011. “Scikit-Learn: Machine Learning in Python.” *Journal of Machine Learning Research* 12 (Oct): 2825–30.

Pérez, F., and B. E. Granger. 2007. “IPython: A System for Interactive Scientific Computing.” *Computing in Science & Engineering* 9 (3): 21–29.

Pérez, F., B. E. Granger, and J. D. Hunter. 2011. “Python: An Ecosystem for Scientific Computing.” *Computing in Science & Engineering* 13 (2): 13–21.

Van der Walt, S., S. C. Colbert, and G. Varoquaux. 2011. “The NumPy Array: A Structure for Efficient Numerical Computation.” *Computing in Science & Engineering* 13 (2): 22–30.

Waskom, M, O Botvinnik, P Hobson, J Warmenhoven, JB Cole, Y Halchenko, J Vanderplas, et al. 2014. *Seaborn: Statistical Data Visualization*. Stanford, California.